SFT:

input comes from human, output comes from LM and the human rates the response.

This information is used for reward modelling (which we can think of as a classifier)

Then, you optimize a policy with that reward modeling.

Self-instruct: synthetic generation

human-in-the-loop: human doing post-processing tasks.

- human preference dataset, human writes input, model outputs, human rates them.
- for the SFT dataset both input and output are written by a human

Distillation:

dSFT (generate multi-turn AI dialogues)
AIF (response generation and AI ranking)
dDPO (distillation of AI preferences)

4 different training techniques

1. Pretraining the LM
2. In context learning (prompt-based learning)
3. Supervised fine-tuning
4. Reinforcement learning with Human Feedback

evaluating a chatbot:

- evaluating instruction following/ chattiness
- evaluating the reward modeling
- red-teaming (or checking response based on adversial prompting)
- helpfulness (elo rating)
- can the model choose between a toxic and an untruthful response? (HF reward model benchmarks)

Red Teaming LLMs:

- there's no leaderboard for this unfortunately.

using llms as an evaluator (gpt4 in particular)

↳ gpt4 prefers gpt inputs to human inputs :)

↳ prefers more wiggle and layer responses

↳ doesn't align with <u>useful</u> relational responses.